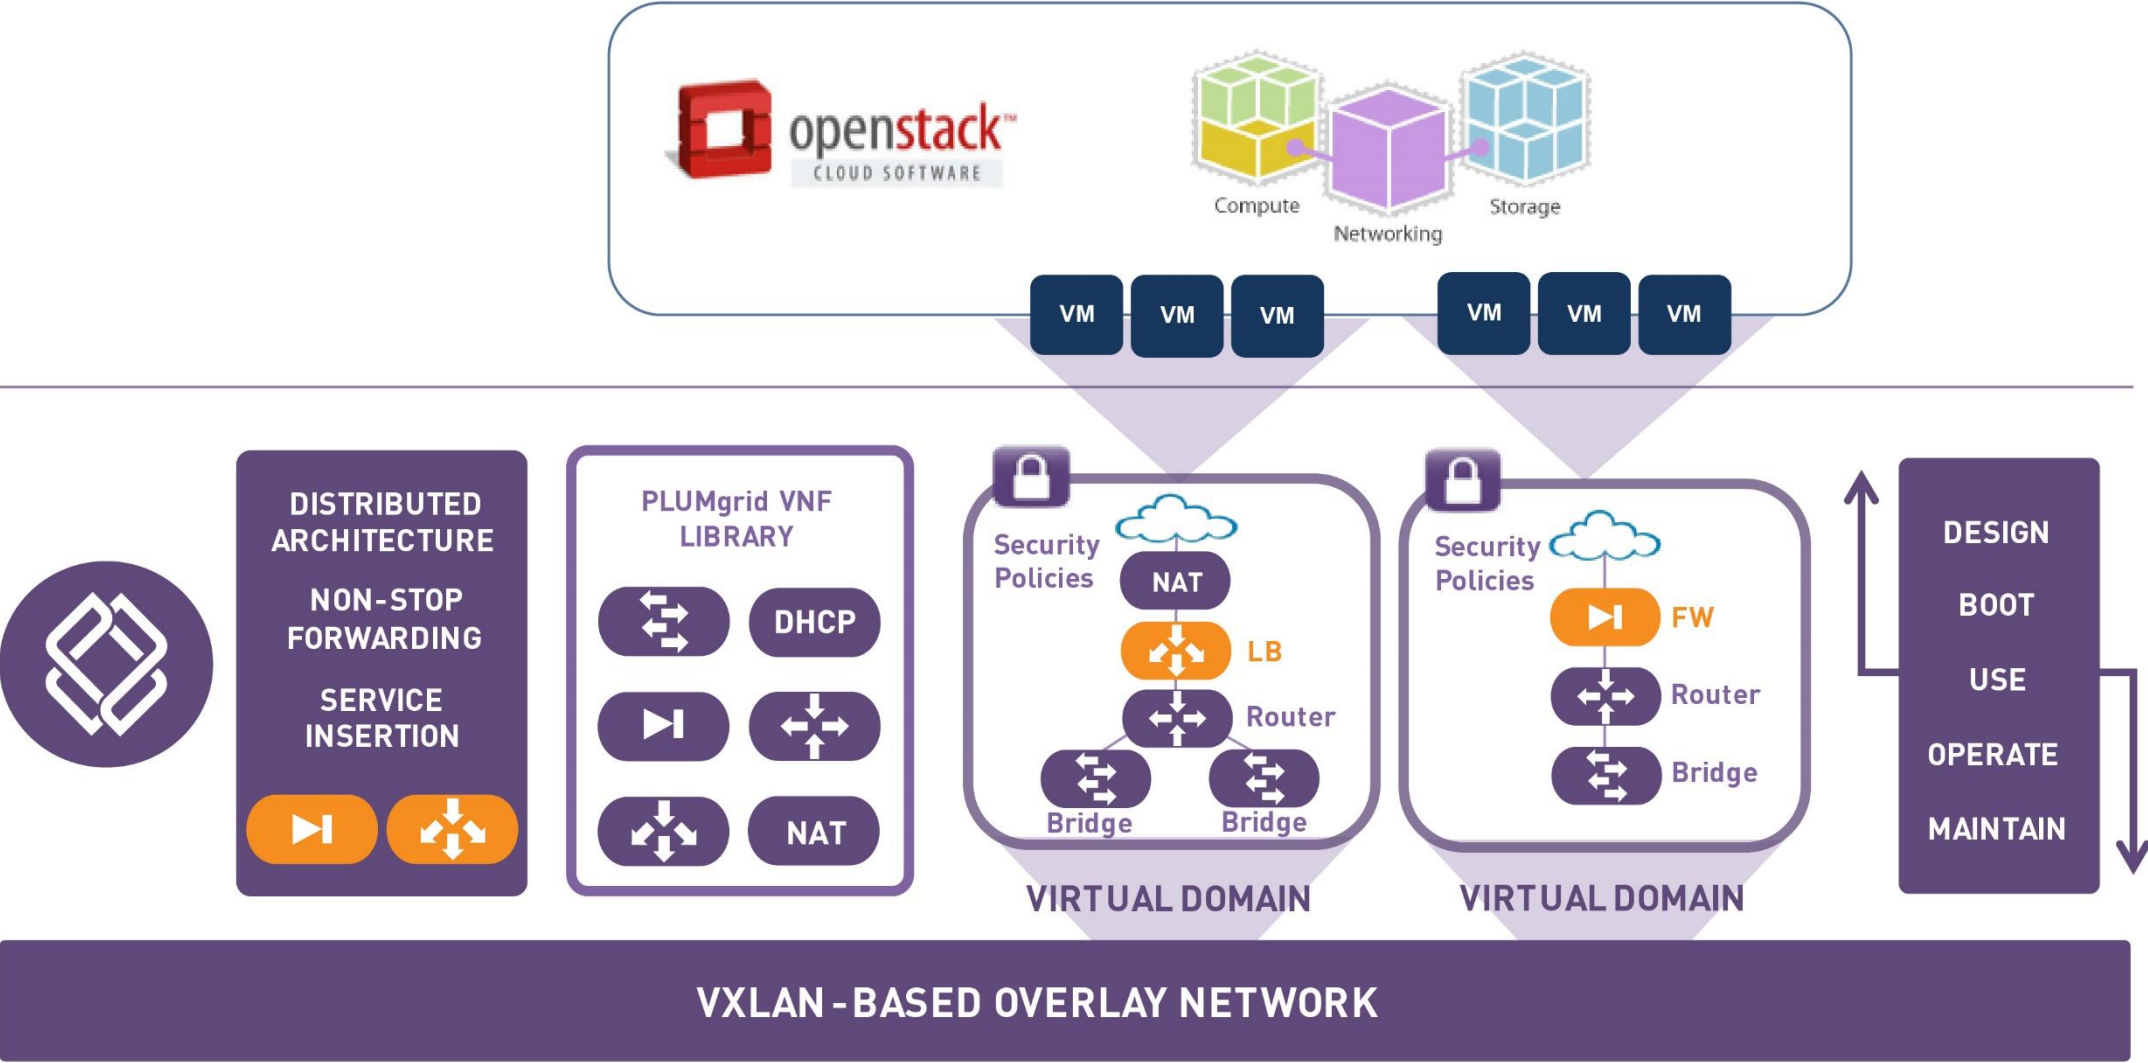


# eBPF and IO Visor: The what, how, and what next!

**Affan A. Syed**

Director Engineering,  
PLUMgrid Inc.

# About PLUMgrid



# PLUMgrid's Mission

*Deliver comprehensive virtual networking solutions  
that scale, secure, and simplify the modern cloud  
data center*

LINUX FOUNDATION COLLABORATIVE PROJECTS



Any  
Hardware

Any  
Network

Any  
Service

Any  
Hypervisor

Any  
Container

Pervasive Scale and Security

# Talk outline

# Outline

- Berkeley Packet Filter (BPF) and bytecte
- Extended BPF (eBPF): Motivation and features
- BCC and IO Visor
- Basic demos
- Research Directions

# Packet Filters and cBPF


# Packet Filters

Objective: To observe **all** traffic but capture only a subset

Problem: Packet traversal through normal stack is slow

Solution: Setup filters **in kernel** where packet dropped if not match

....but these filters need to be secure!



# Tcpdump our friend

Lets do: `sudo tcpdump -p -ni eth0 "ip and udp"`

Now lets do: `$ sudo tcpdump -p -ni eth0 -d "ip and udp"`

```
(000) ldh    [12]
(001) jeq    #0x800      jt 2   jf 5
(002) ldb    [23]
(003) jeq    #0x11      jt 4   jf 5
(004) ret    #65535
(005) ret    #0
```

This code runs for every packet that arrives on eth0



# The concept of pseudo-machine

Think Java --- but don't think VM !


Virtualize a machine instruction set

Write byte code for this “fictional” machine

***verify** this code is loop free and optimized*

Interpret, in-kernel, for *\*any\** real processor

# BPF Overview (insertion and usage)



## Other uses and extension

Slowly evolving cBPF

seccomp support for sandboxing

*tc* filter for traffic shaping

JIT compiler

# Extending BPF

... while building a programmable Data Plane

# A new SDN architecture

operators



Management API



Closed Network Functions

*Closed* North Bound API



Controller

South Bound API, *but no extensibility*



Data Plane

**Traditional  
SDN architecture**

developers




operators



Management API

Controller



CP

CP

CP

Network Function

CP-DP APIs


DP

DP

DP

IO Visor

SDK



developers

**With a programmable DP and  
controller on the side**

# Goals for a programmable Data Plane = eBPF

Enable packet parsing, lookup, modification, and updates

Guarantee safety of code run on a production system

Native performance

# eBPF as a syscall interface

Introduced as a separate syscall (user space access)

```
int bpf(int cmd, union bpf_attr *attr, unsigned int size);
```

linux 3.15 and above

Moved out of the networking subsystem

Streamlined use, extensibility through single API

# Enhanced architecture

| classic BPF   | extended BPF  |
|---|---|
| 2 registers + stack<br>32-bit registers<br>4-byte load/store to stack<br>1-4 byte load from packet<br>Conditional jump forward<br>+, -, *, ... instructions | <b>10 registers</b> + stack<br><b>64-bit registers</b> with 32-bit sub-registers<br>1-8 byte load/store to stack, maps, context<br>Same + store to packet<br>Conditional jump forward and backward<br>Same + signed_shift + endian<br><b>Call instruction</b><br><b>tail_call</b><br><b>map lookup/update/delete helpers</b><br><b>packet rewrite, csum, clone_redirect</b><br>sk_buff read/write |

- Can build more complicated program
- Faster interpretation and JIT
- Support for calls to *approved* helper functions



# Maps

Maps = <key, value> storage

Save state across invocation of programs in kernel = state machines!

Example: *fd bpf\_table\_lookup(table\_id, key)*

Userspace can create/access/delete these maps (using bpf syscall)

*loosely coupled communication between user and kernel space*

Maps can be shared between eBPF programs

*HASH, ARRAY ... and growing*

Stateful programmability and async interaction with user space

# Helper functions and tail calls

Invoke sanitized functions from within the eBPF program

*like a library – but ... of course .. In-kernel*

*e.g u64 bpf\_ktime\_get\_ns(void), int bpf\_trace\_printk(const char \*fmt, int fmt\_size, ...), u32 prandom\_u32(void)*

Tail call feature a combo of two components

*bpf\_tail\_call(ctx, prog\_array\_map, index)*

*and PROG\_ARRAY\_MAP*

increased capabilities, and sanitized access

# Summary (for later reference)

|                        |   |   |
|------------------------|---|---|
| eBPF<br>maps           | BPF_MAP_TYPE_HASH                             | Optimized for speed of lookup and atomic updates  |
|                        | BPF_MAP_TYPE_ARRAY                            | Fixed (4 byte) key to index into the array, thus giving fastest possible lookup. Array elements are zero initialized.   |
|                        | BPF_MAP_TYPE_PROG_ARRAY                       | Like Array Maps, but value also of only 4 bytes representing file descriptors referring to other eBPF programs.   |
|                        | BPF_MAP_TYPE_PERF_EVENT_ARRAY                 | Like Array Maps, but value containing pointers referring to kernel perf events.   |
| eBPF<br>map<br>helpers | BPF_MAP_{CREATE, LOOKUP, UPDATE, DELETE}_ELEM | In order to create/lookup/update(create or update)/delete elements in the maps.   |
|                        | BPF_MAP_GET_NEXT_KEY                          | Looks up an element by key in the map referred to by the file descriptor fd and sets the next_key pointer to the key of the next element.                               |
|                        | close   | Delete a map  |
|                        | BPF_PROG_TYPE_SOCKET_FILTER                   | Attach an eBPF program when you create a socket (tcp, udp, raw, unix, etc.)   |
| eBPF<br>programs       | BPF_PROG_TYPE_KPROBE                          | Attach an eBPF program to a kprobe events (particular kernel function), triggers when the kernel function is called, giving users dynamic visibility inside the kernel. |
|                        | BPF_PROG_TYPE_SCHED_CLS                       | Attach an eBPF program to Linux TC classifier.  |
|                        | BPF_PROG_TYPE_SCHED_ACT                       | Attach an eBPF program to Linux TC action.  |

# Verifier and kernel safety

eBPF new architecture more complex

*required a **brand** new verifier*

Provably confirms inserted program **does not:**

*create loops, delays execution interminably, illegally  
dereference pointers*

Done statically, one-time, with an exhaustive search

*some heuristics to improve verification time*



# LLVM compiler, Interpreter, and JIT

Restricted “C” code that compiles to bpf bytecode

*LLVM backend for this purpose with clang frontend*

Once inserted, the code is “hooked” to a kernel event

*no sense hooking to userspace events!*

On event firing the appropriate code is run in either

*native or interpreted mode*


# Example Hooks: The networking stack

traffic control (TC):

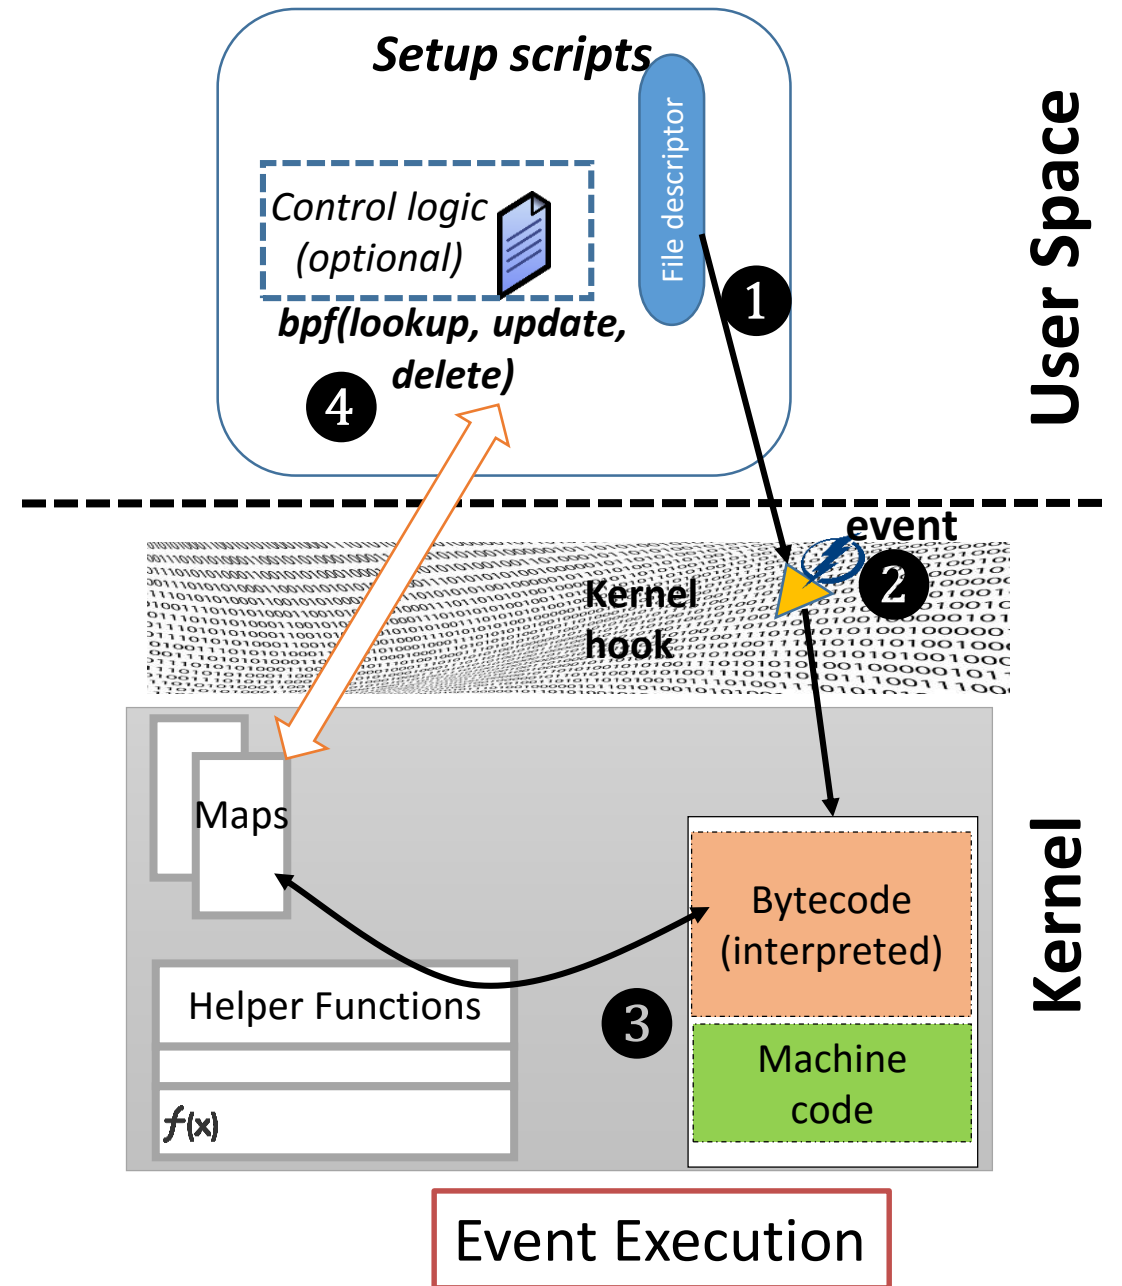
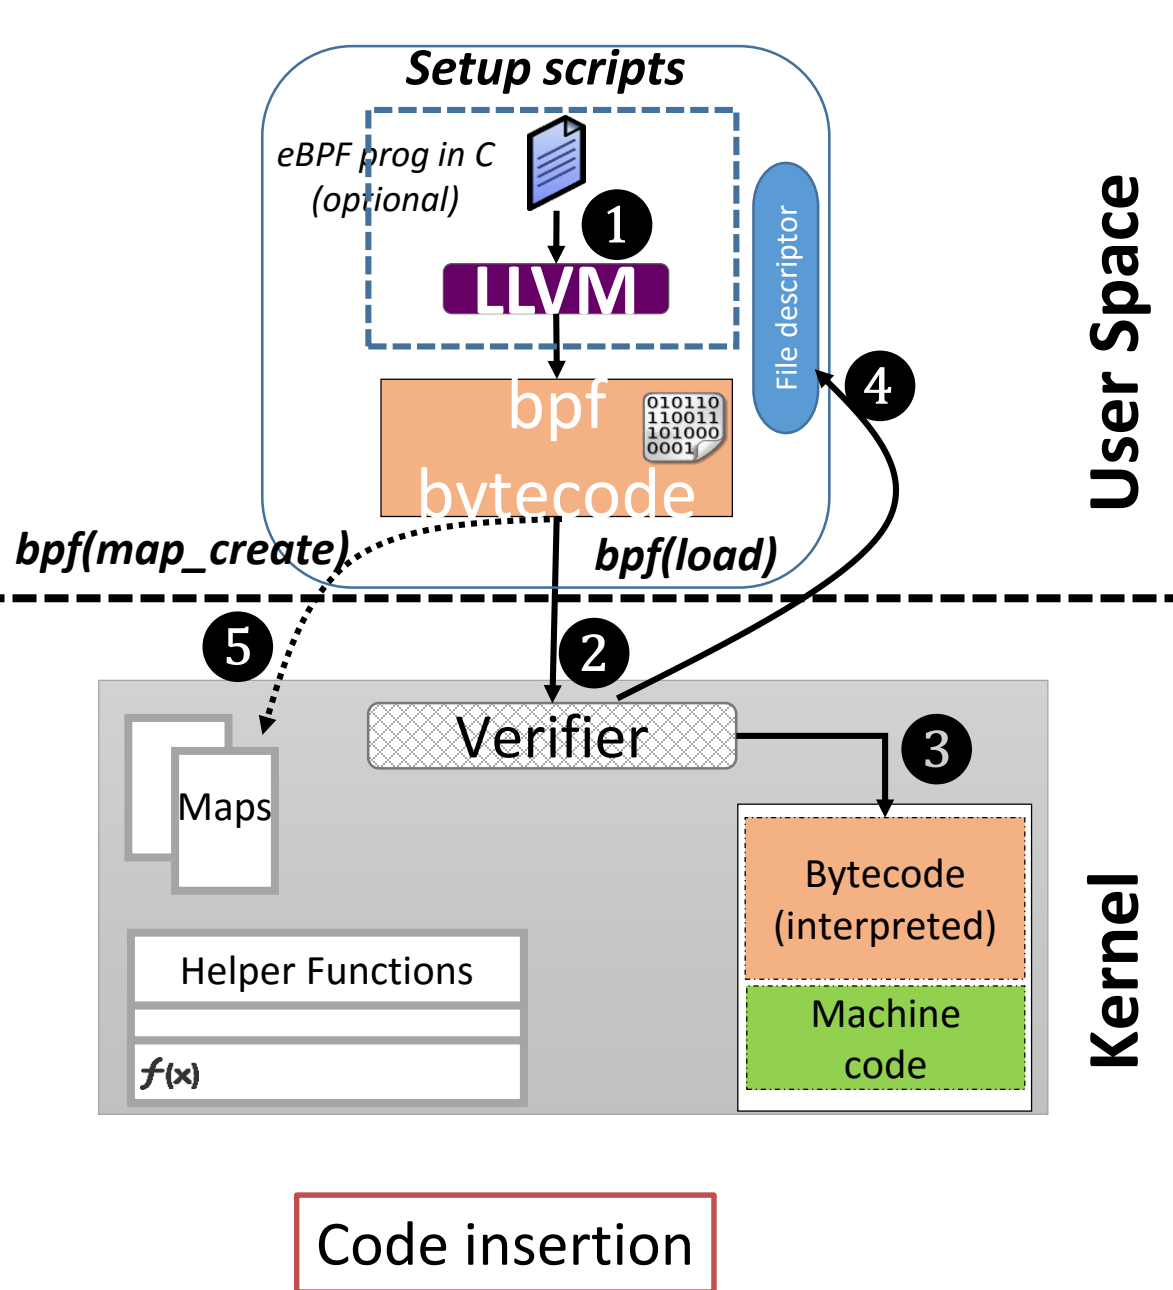
queues (classification or action time)

sockets: STREAM  
(L4/UDP), DATAGRAM  
(L4/TCP) or RAW

others: kprobes, syscalls,  
tracepoints ...



# Visual Flow of code insertion and use



**eBPF, IO Visor and BCC**



# Complexity of making eBPF code


Writing eBPF programs was “brutal”

Even with compiler, the map/code sharing

Enter BPF compiler collection (BCC)

## The Tracing Landscape, ~~May~~ 2015

Sep



Make using features of eBPF easier to use

Python front-end and scripts to

*create/access/delete maps*

*load programs from a restricted “C” format*

*attach to different locations with a simple API*

# BCC and a few screen-shots!



## BPF Compi

BCC is a toolkit for creating examples. It makes use of € of what BCC uses requires

## Kernel requirements

### Requirements

In general, to use these features, a Linux kernel version 4.1 or newer is required. In addition, the following flags should be set:

```
CONFIG_BPF=y
CONFIG_BPF_SYSCALL=y
# [optional, for tc filters]
CONFIG_NET_CLS_BPF=m
# [optional, for tc actions]
CONFIG_NET_ACT_BPF=m
CONFIG_BPF_JIT=y
CONFIG_HAVE_BPF_JIT=y
# [optional, for kprobes]
CONFIG_BPF_EVENTS=y
```


# IO Visor Project

Linux Foundation Collaborative Project


**IO Visor Engine** is an abstraction of an IO execution engine

A set of development tools, **IO Visor Dev and Management**

A set of **use cases & applications** like Networking, Security, Tracing & others

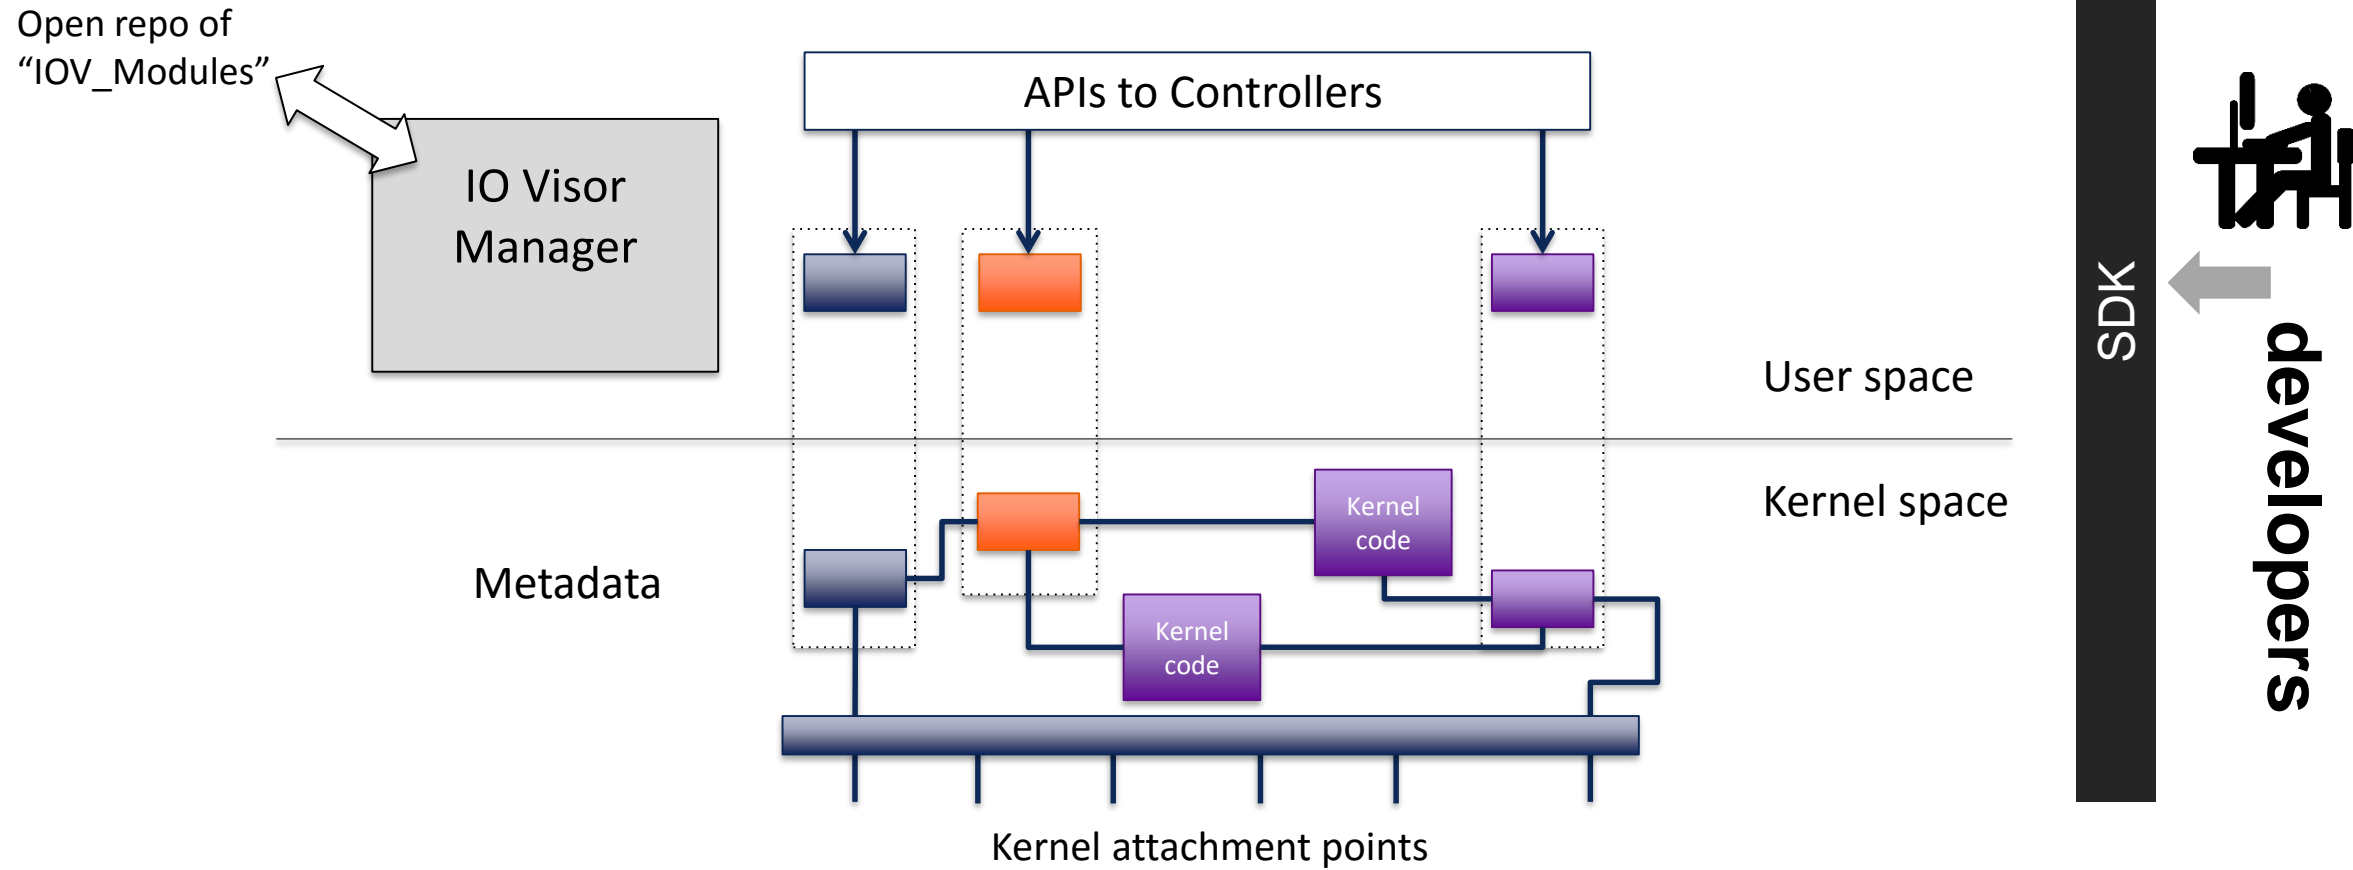


# Founding Members



[www.iovisor.org](http://www.iovisor.org)

# IO Visor Modules



<https://github.com/iovisor/>

Lots of interesting projects there right now  
bpf-based file system (FUSE)

Join and contribute!

# Useful links

- <https://github.com/iovisor/bcc>
- <https://github.com/iovisor/bpf-docs>
- <http://lwn.net/Articles/603984/>
- <http://lwn.net/Articles/603983/>
- <https://lwn.net/Articles/625224/>
- <https://www.kernel.org/doc/Documentation/networking/filter.txt>
- <http://man7.org/linux/man-pages/man2/bpf.2.html>
- [https://linuxplumbersconf.org/2015/ocw//system/presentations/3249/original/bpf\\_llvm\\_2015aug19.pdf](https://linuxplumbersconf.org/2015/ocw//system/presentations/3249/original/bpf_llvm_2015aug19.pdf)
- [https://videos.cdn.redhat.com/summit2015/presentations/13737\\_an-overview-of-linux-networking-subsystem-extended-bpf.pdf](https://videos.cdn.redhat.com/summit2015/presentations/13737_an-overview-of-linux-networking-subsystem-extended-bpf.pdf)
- <https://github.com/torvalds/linux/tree/master/samples/bpf>
- <http://events.linuxfoundation.org/sites/events/files/slides/tracing-linux-ezannoni->
- [https://www.kernel.org/doc/Documentation/prctl/seccomp\\_filter.txt](https://www.kernel.org/doc/Documentation/prctl/seccomp_filter.txt)
- [http://lxr.free-electrons.com/source/net/sched/cls\\_bpf.c](http://lxr.free-electrons.com/source/net/sched/cls_bpf.c)



# Live Demo

... and prayers!

# Things we demo

## Hello world

*attach to clone event, print hello every time*

## Hello world and state

*sum the number of events*

## Networking example

*more complicated but fun 😊*

# Research Threads

... for the adventurous amongst you!

# Networking and packet manipulations

Increasing application response time

frequent /cached responses in kernel

Container Networking

custom encapsulation protocols and metadata

State-full QoS


Fast, flexible and state-full Firewalls

System call trapping and *dynamic* taint analysis

*Faster and less resource hungry analysis*

## Disk latency heat-maps

<https://github.com/torvalds/linux/tree/master/samples/bpf/>



One map for time-stamp, other for latency

Event-based, packet-based micro-kernel abstraction

*Saving energy with operation in-kernel,  
user-space tools for config/mgmt/debug  
think tinyOS-inside-Linux*

Software defining an API for heterogenous and  
opportunistic low-power coms

*wifi, zigbee, Z-wave*

Thank you!  
Questions?

*We are hiring*  
[hr.isb@plumgrid.com](mailto:hr.isb@plumgrid.com)